


Edge Computing of YOLO for use in Fall Detection

Ben Hempelmann and Cameron Legg

Project Overview

Original Plan


-
- Detect human falls using a computer vision approach
 - This does not require wearables or other sensors
 - Run YOLO Darknet on a Raspberry Pi, using the OAK-D camera and its built-in hardware accelerator and compare with Coral USB TPU



Changed Plans

- Run YOLO Darknet on Nvidia Jetson Orin NX and Oak-D Camera
 - Compute inferences on
 - Oak-D Camera
 - GPU
 - CPU
 - For benchmarking a 5 minute video was passed as a stream to the models
 - The camera feature of the OAK-D was not used, just the processor
 - The benchmarking scripts took the average power consumption and fps every 300 frames.

YOLO



- **Convolutional Neural Network**
- How Fall Detection Works
 - 3 Classes: Sitting, Walking, Fall_Detected
 - Fall_Detected is triggered when someone is laying down
- Parallelization Improves Performance
 - Based on CNNs
 - Applies kernels across different parts of the image
 - Many matrix operations




Training YOLO

- Found a dataset of Labeled Walking/Sitting/Fall_Detected images
 - 374 Images for training / 111 for validation
- Trained a YOLOv11 model using Google Colab
- Model trained on 640p images

Acceleration Frameworks


Nvidia Jetson Orin NX

-
- Ampere GPU – 1024 Nvidia Cuda Cores
 - 8 Core Arm Cortex-A78AE CPU



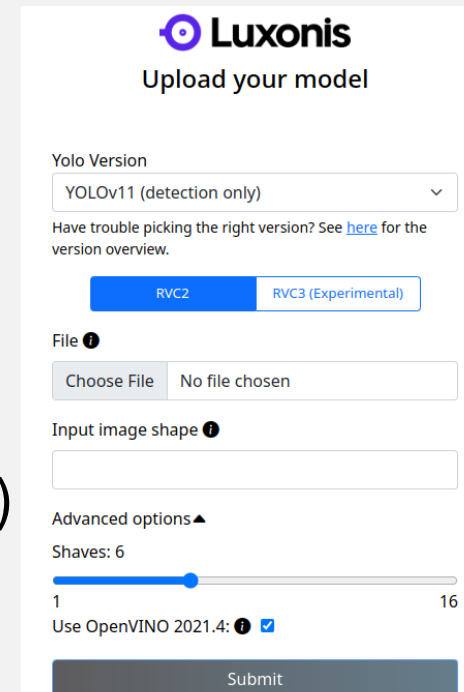
OAK-D Camera

- Based on Robotics Vision Core 2 (RVC2)
 - System on a Chip Architecture
 - Uses Intel Movidius Myriad X Vision Processing Unit
 - Used for on-device deep neural network and computer vision applications
 - 2xLeon CPU cores
 - 16x SHAVE cores
 - Vector Processing Units
 - Max power consumption of 4.5W
 - 2x NCEs (Neural Compute Engines)



Runtime Framework / Programming Interface

- Used Depth AI Python Library
- Runs the YOLO model after it has been converted to OpenVino format
- OpenVino required for running on Movidius Myriad X
- Oak-Models are compiled with number of shaves in mind
- Difficult part was converting the model
- CPU/GPU
 - PyTorch Framework – CUDA and OpenMP (Via Ultralytics library)



The screenshot shows the Luxonis web interface for uploading a model. At the top, the Luxonis logo is displayed with the text "Upload your model". Below this, there is a "Yolo Version" dropdown menu currently set to "YOLOv11 (detection only)". A link is provided for users having trouble picking the right version. Two buttons, "RVC2" and "RVC3 (Experimental)", are visible. The "File" section includes a "Choose File" button and a "No file chosen" status. The "Input image shape" field is empty. Under "Advanced options", the "Shaves" slider is set to 6, with a range from 1 to 16. A checkbox for "Use OpenVINO 2021.4" is checked. A "Submit" button is at the bottom.

Luxonis
Upload your model

Yolo Version
YOLOv11 (detection only) ▼

Have trouble picking the right version? See [here](#) for the version overview.

RVC2 RVC3 (Experimental)

File ⓘ
Choose File No file chosen

Input image shape ⓘ

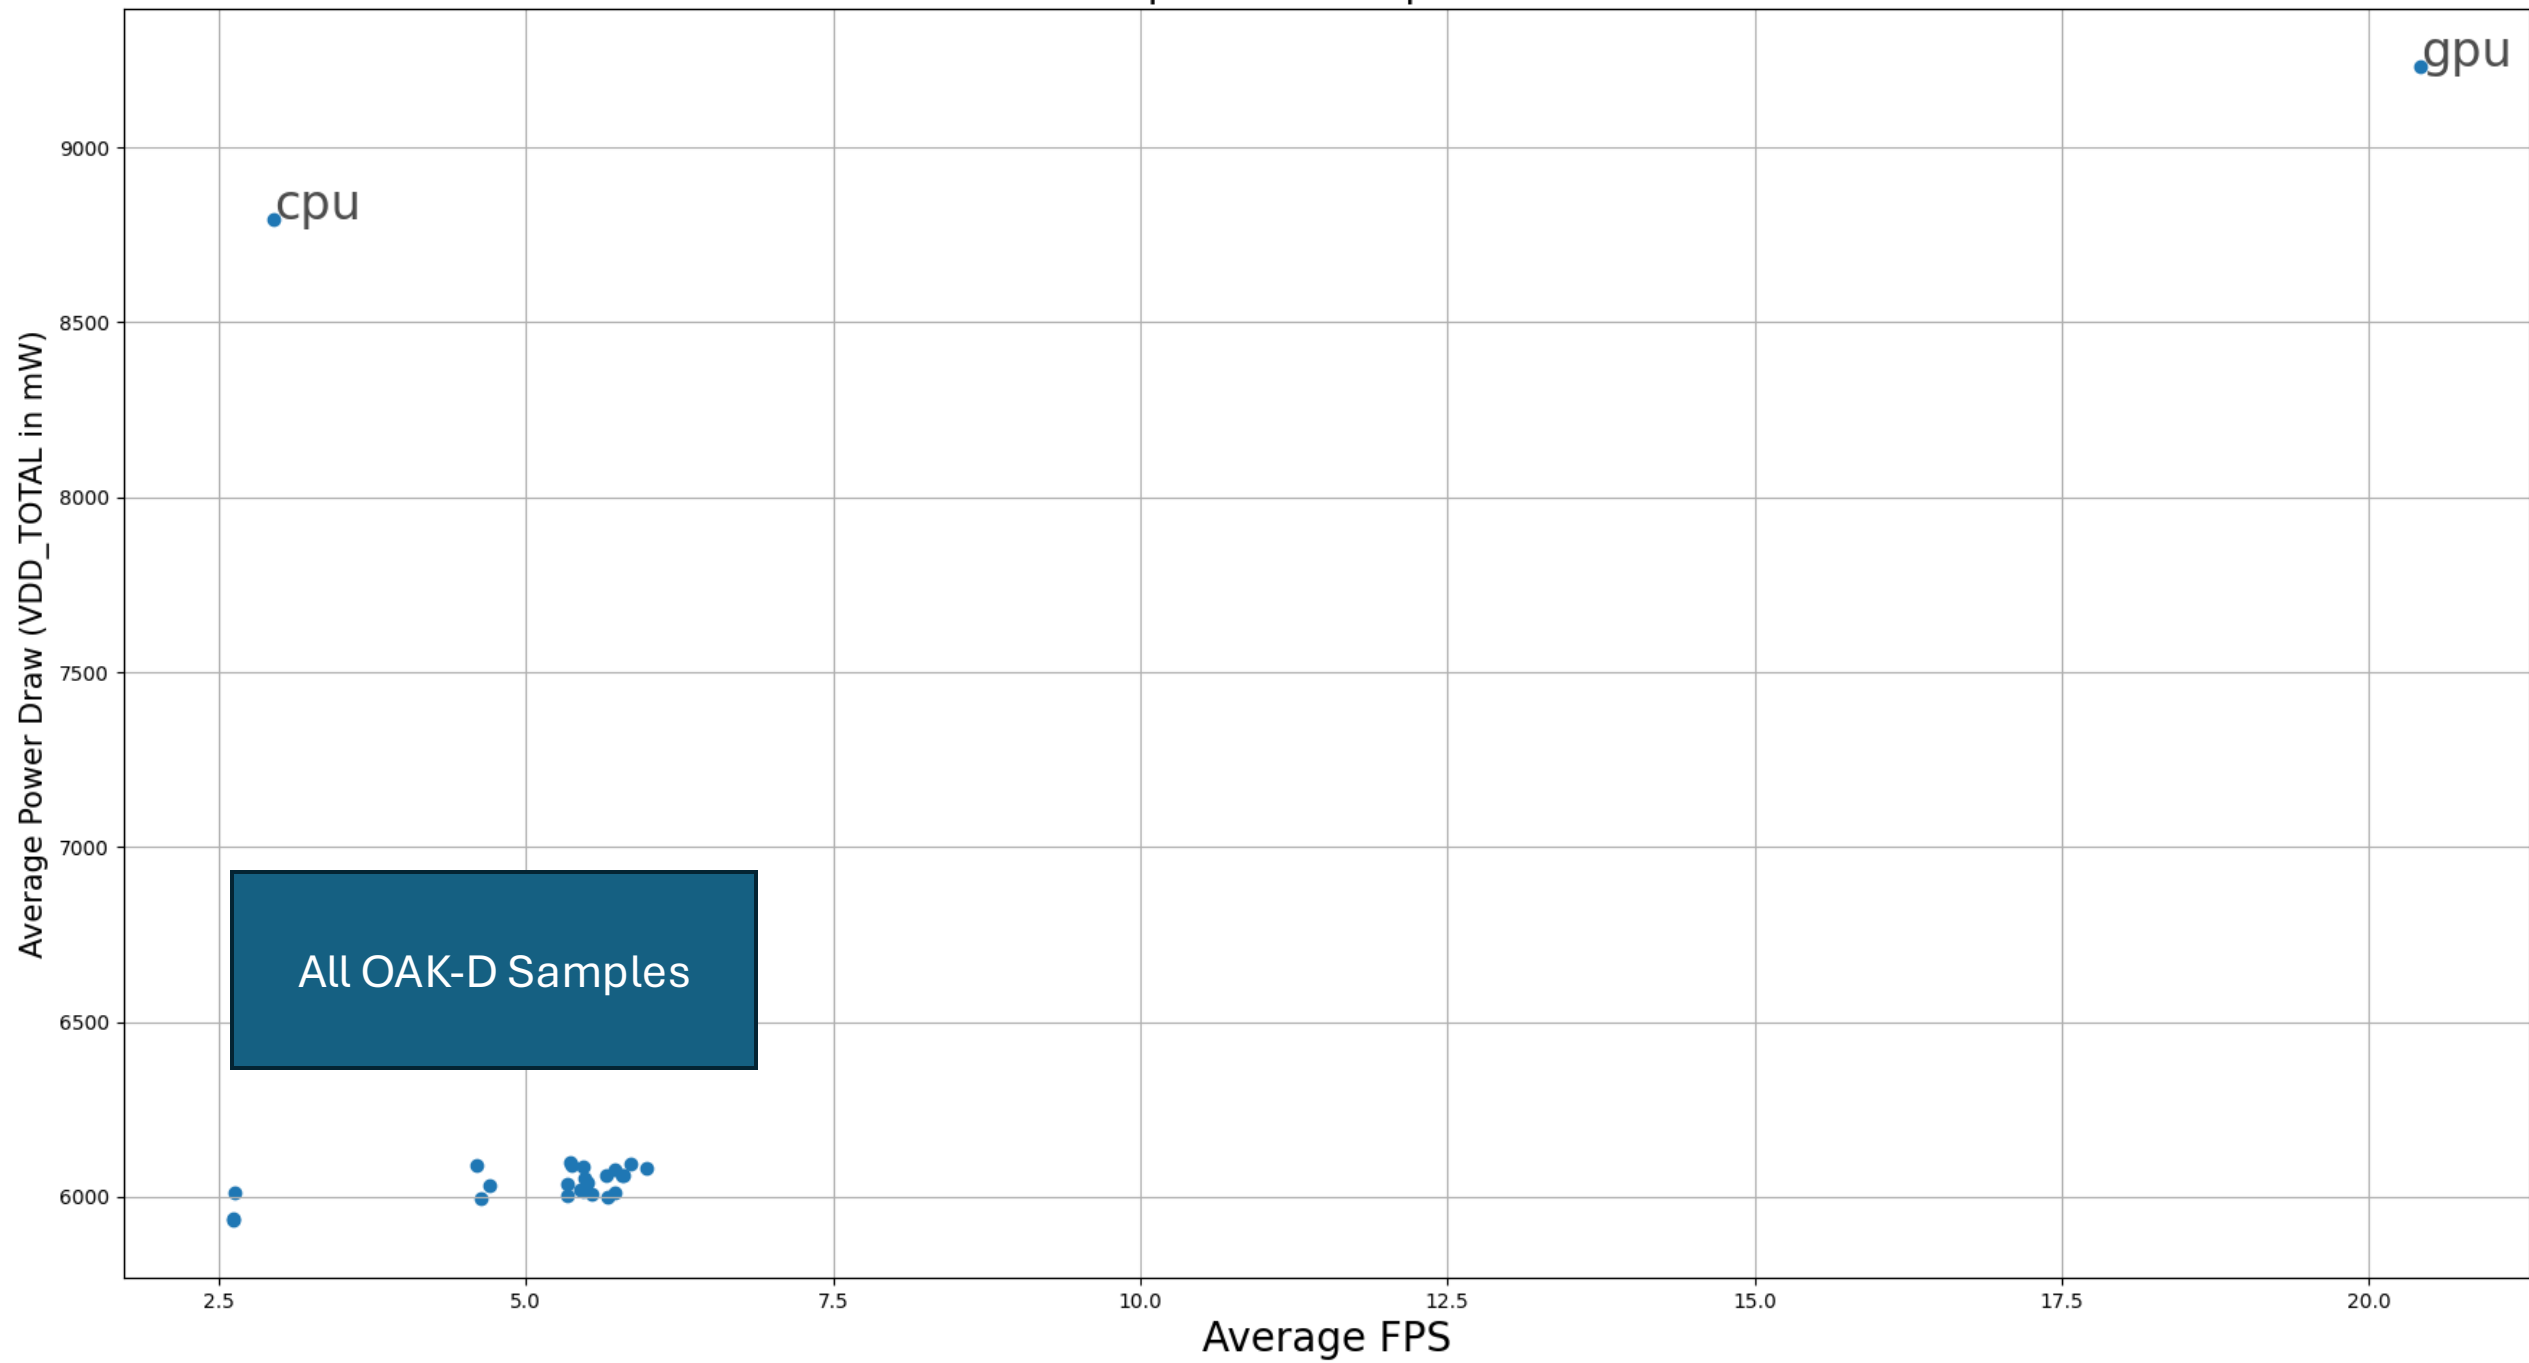
Advanced options ▲
Shaves: 6
1 16

Use OpenVINO 2021.4: ⓘ ☒

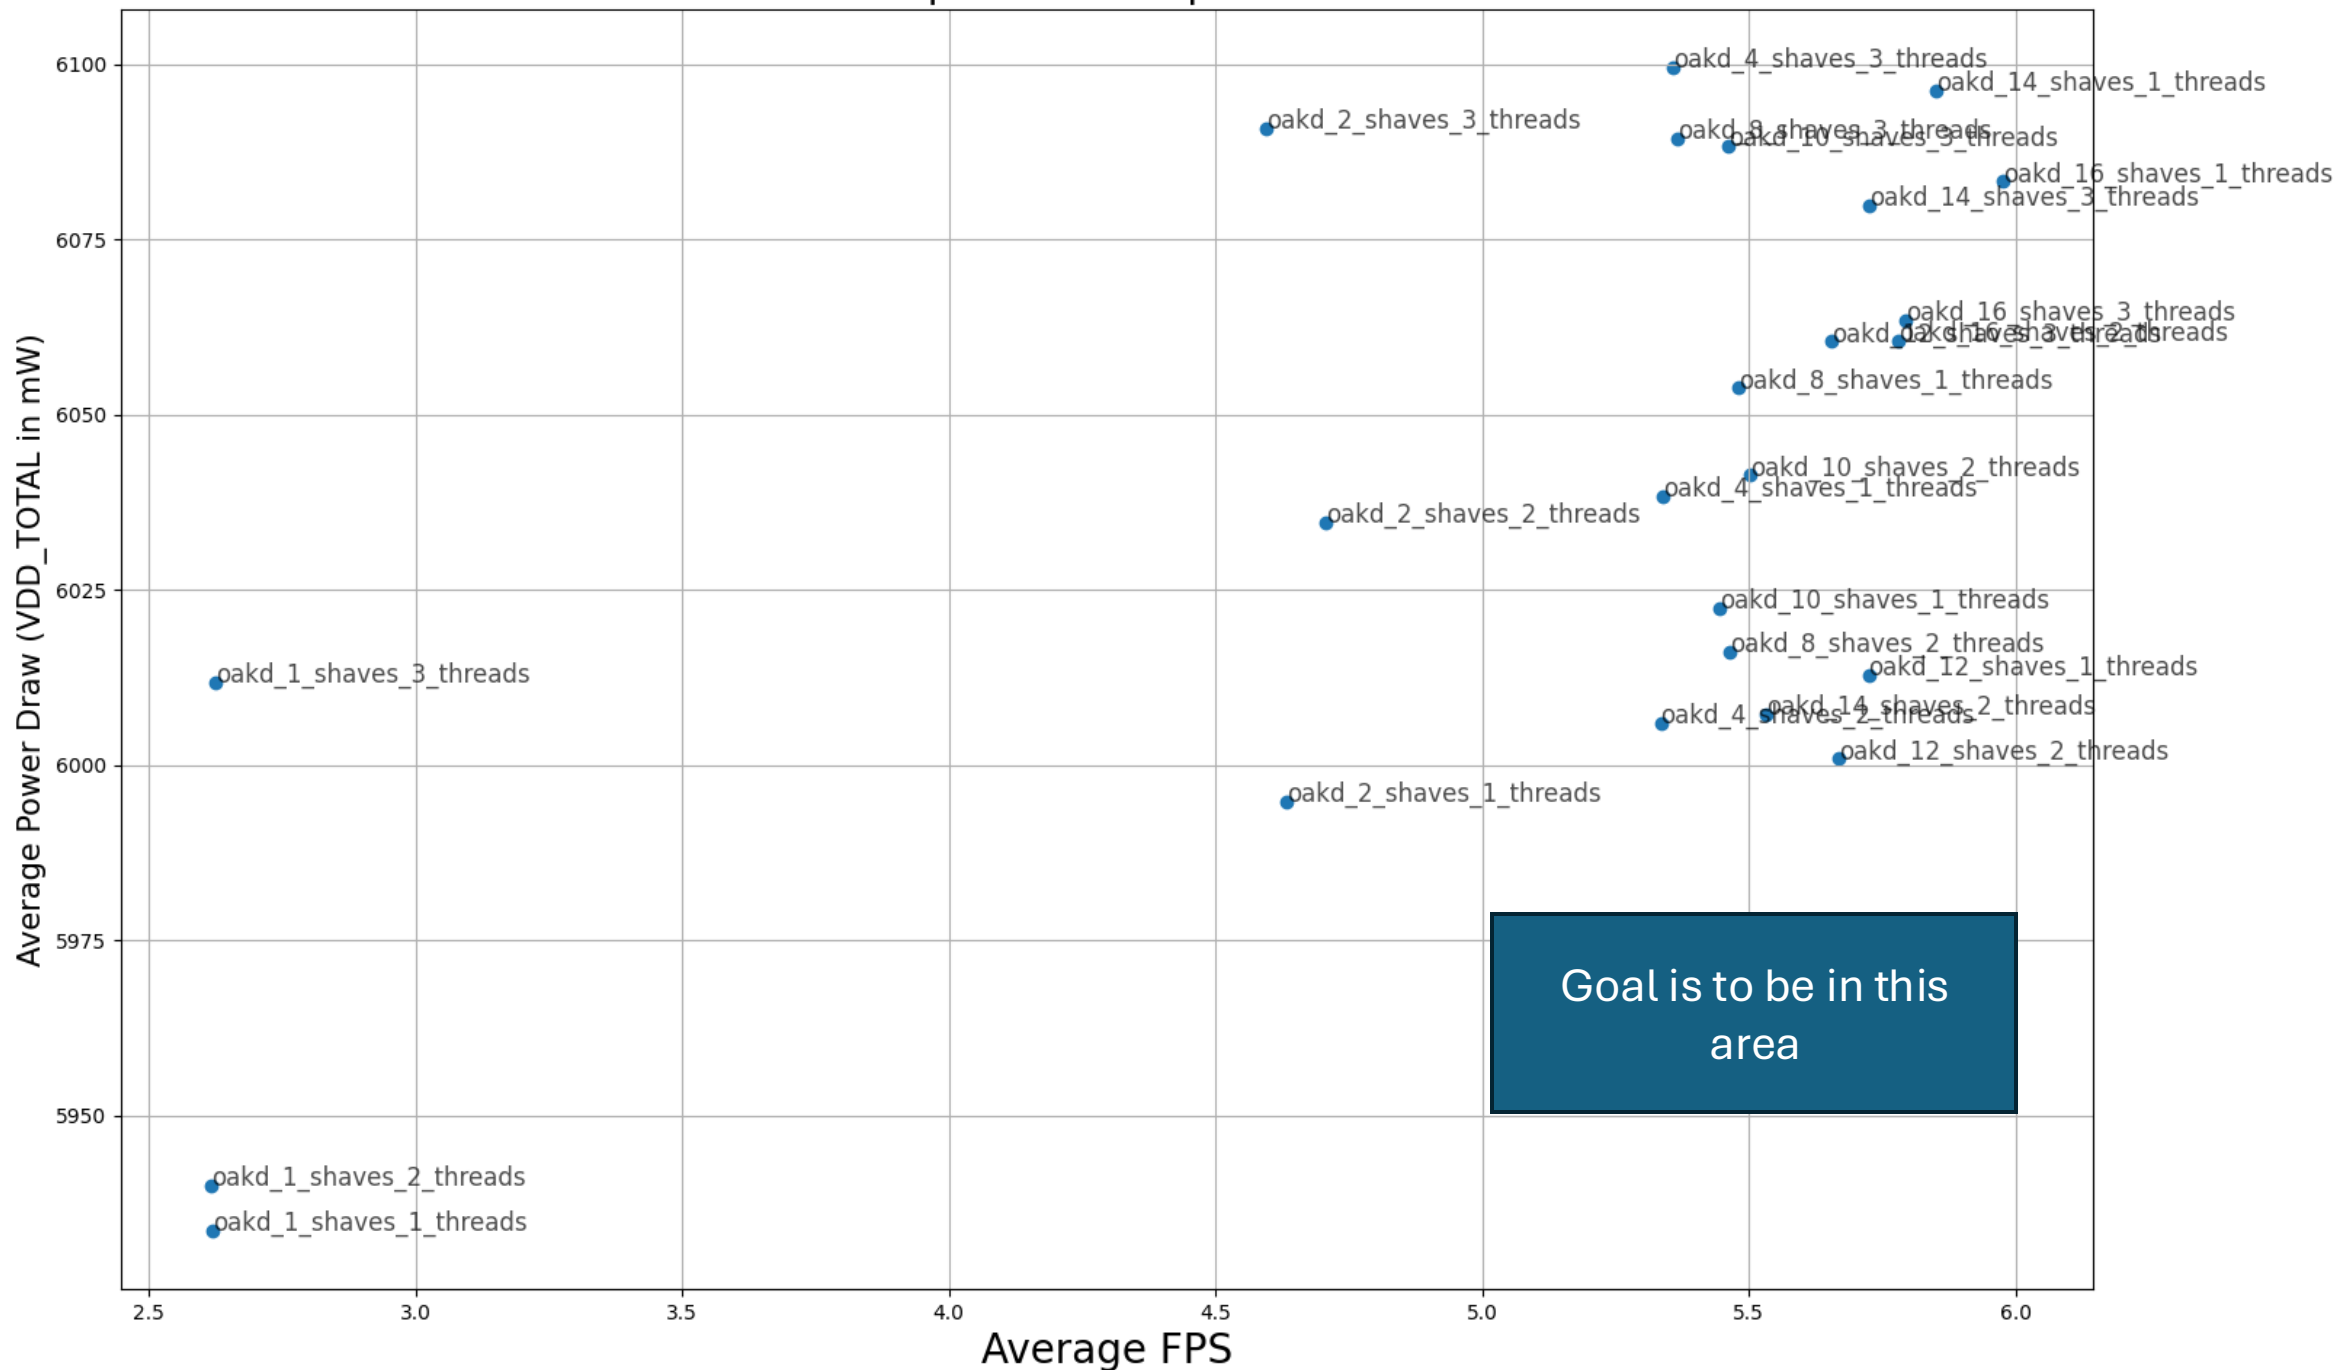
Submit

Performance Comparisons


Power Consumption vs FPS per Benchmark Run



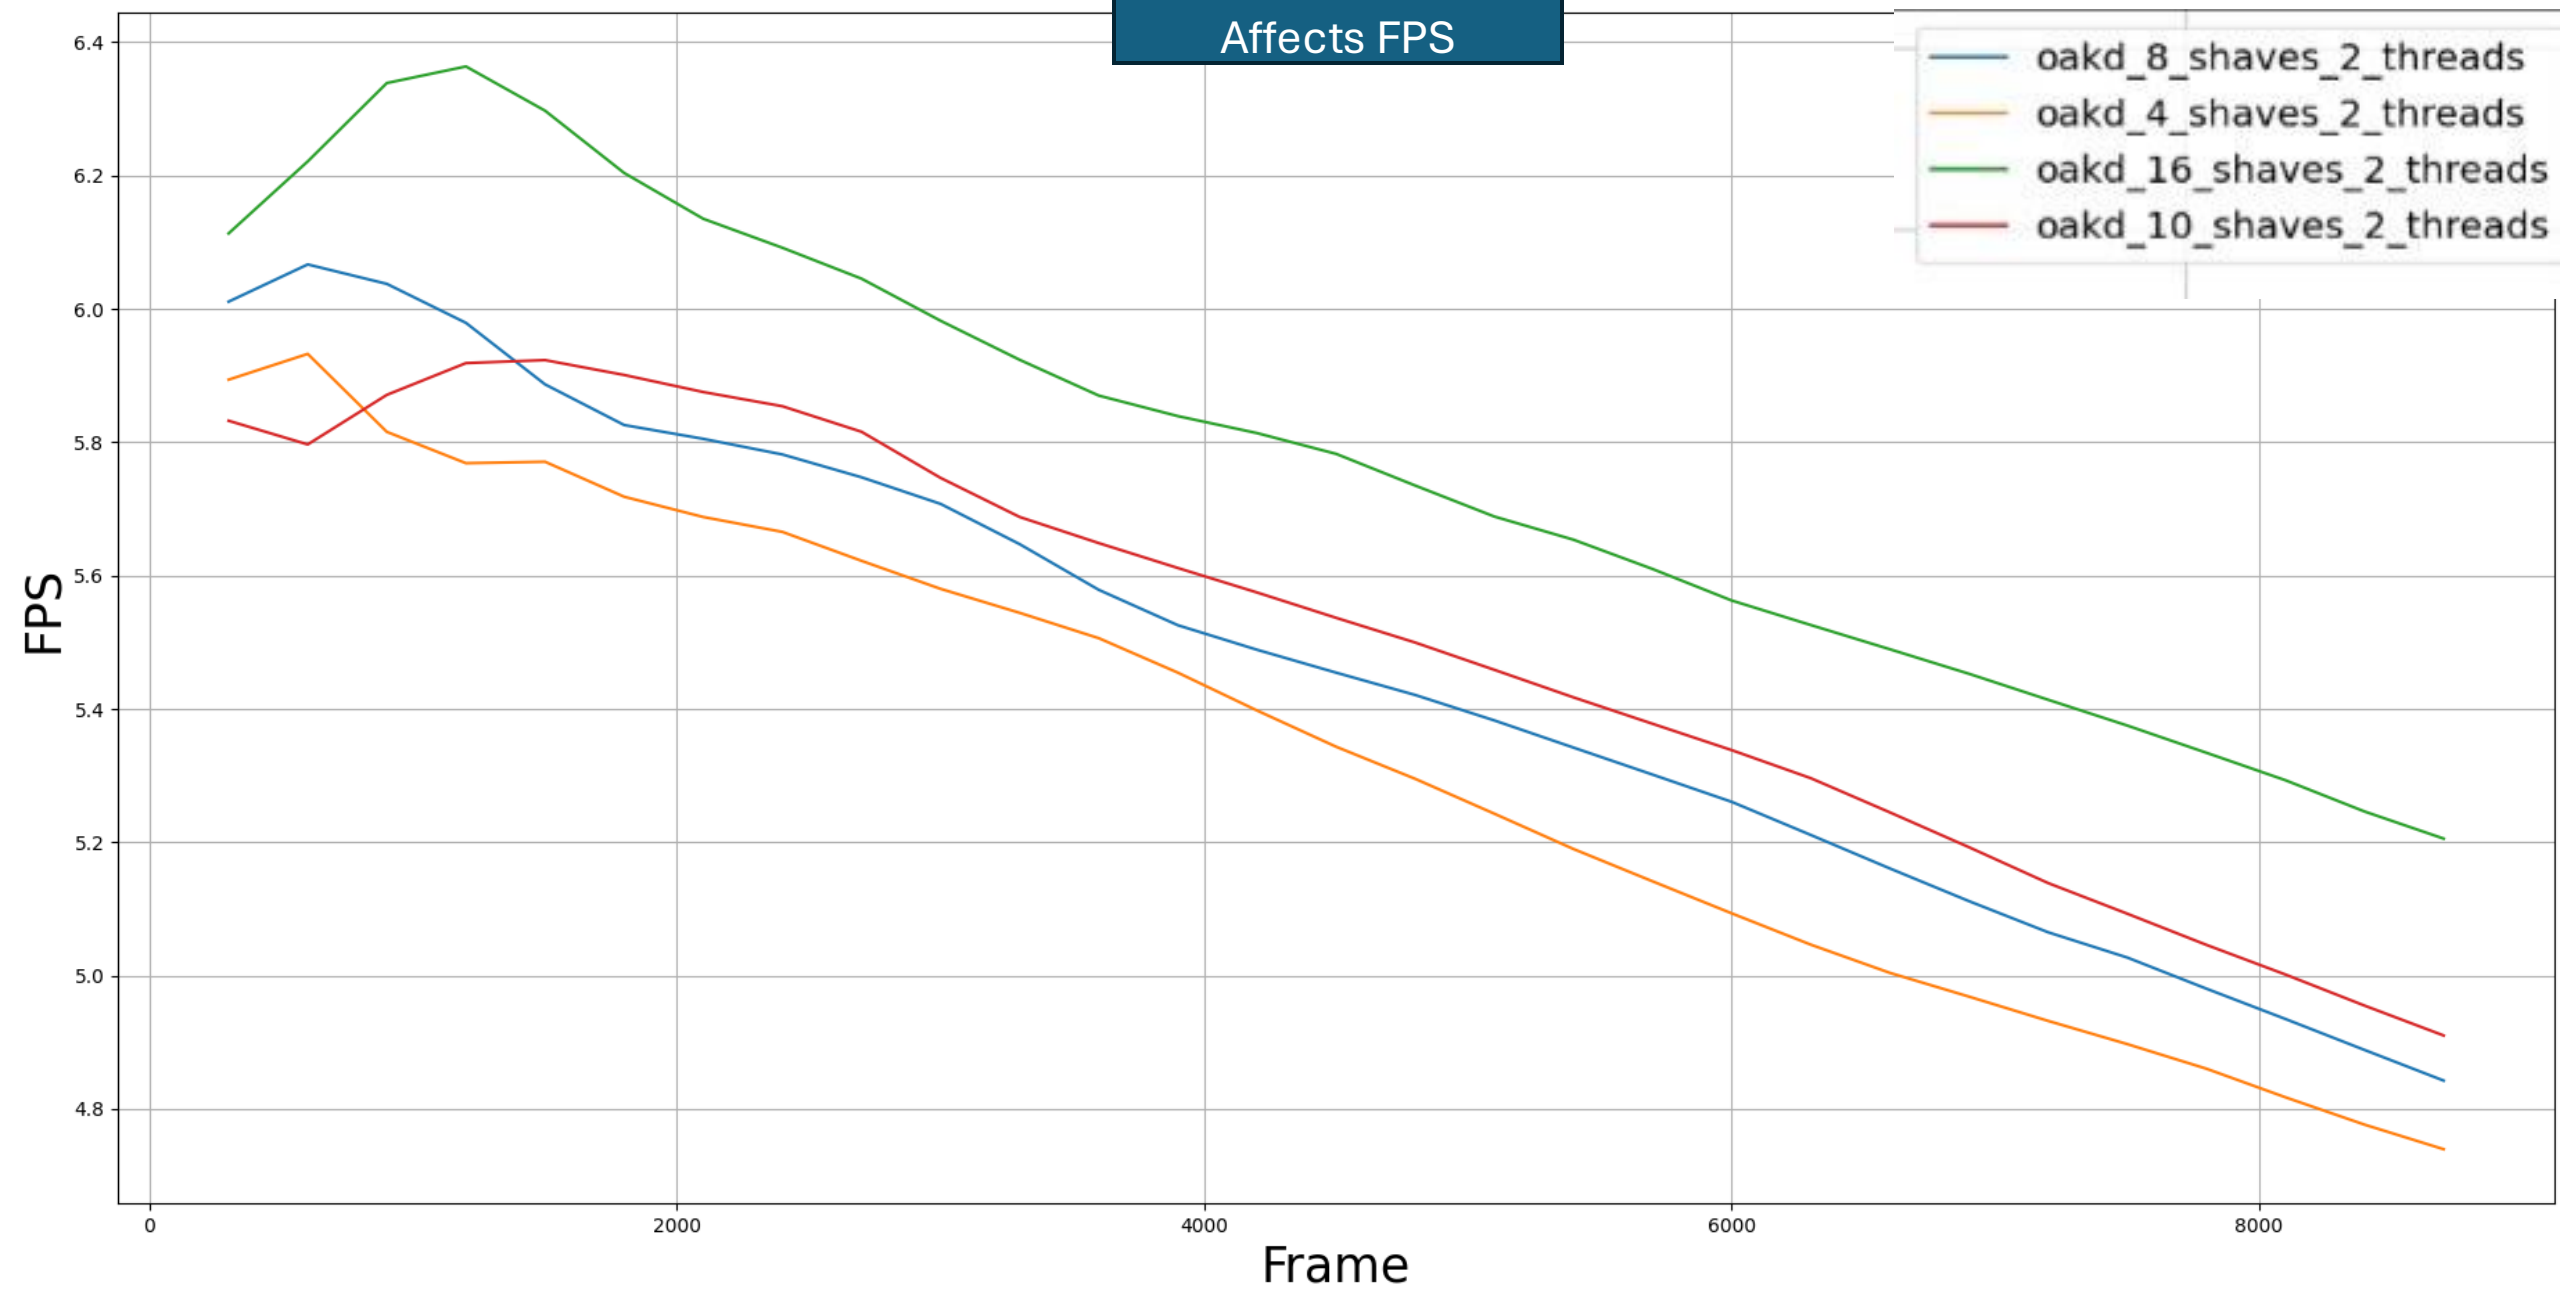
Power Consumption vs FPS per Benchmark Run



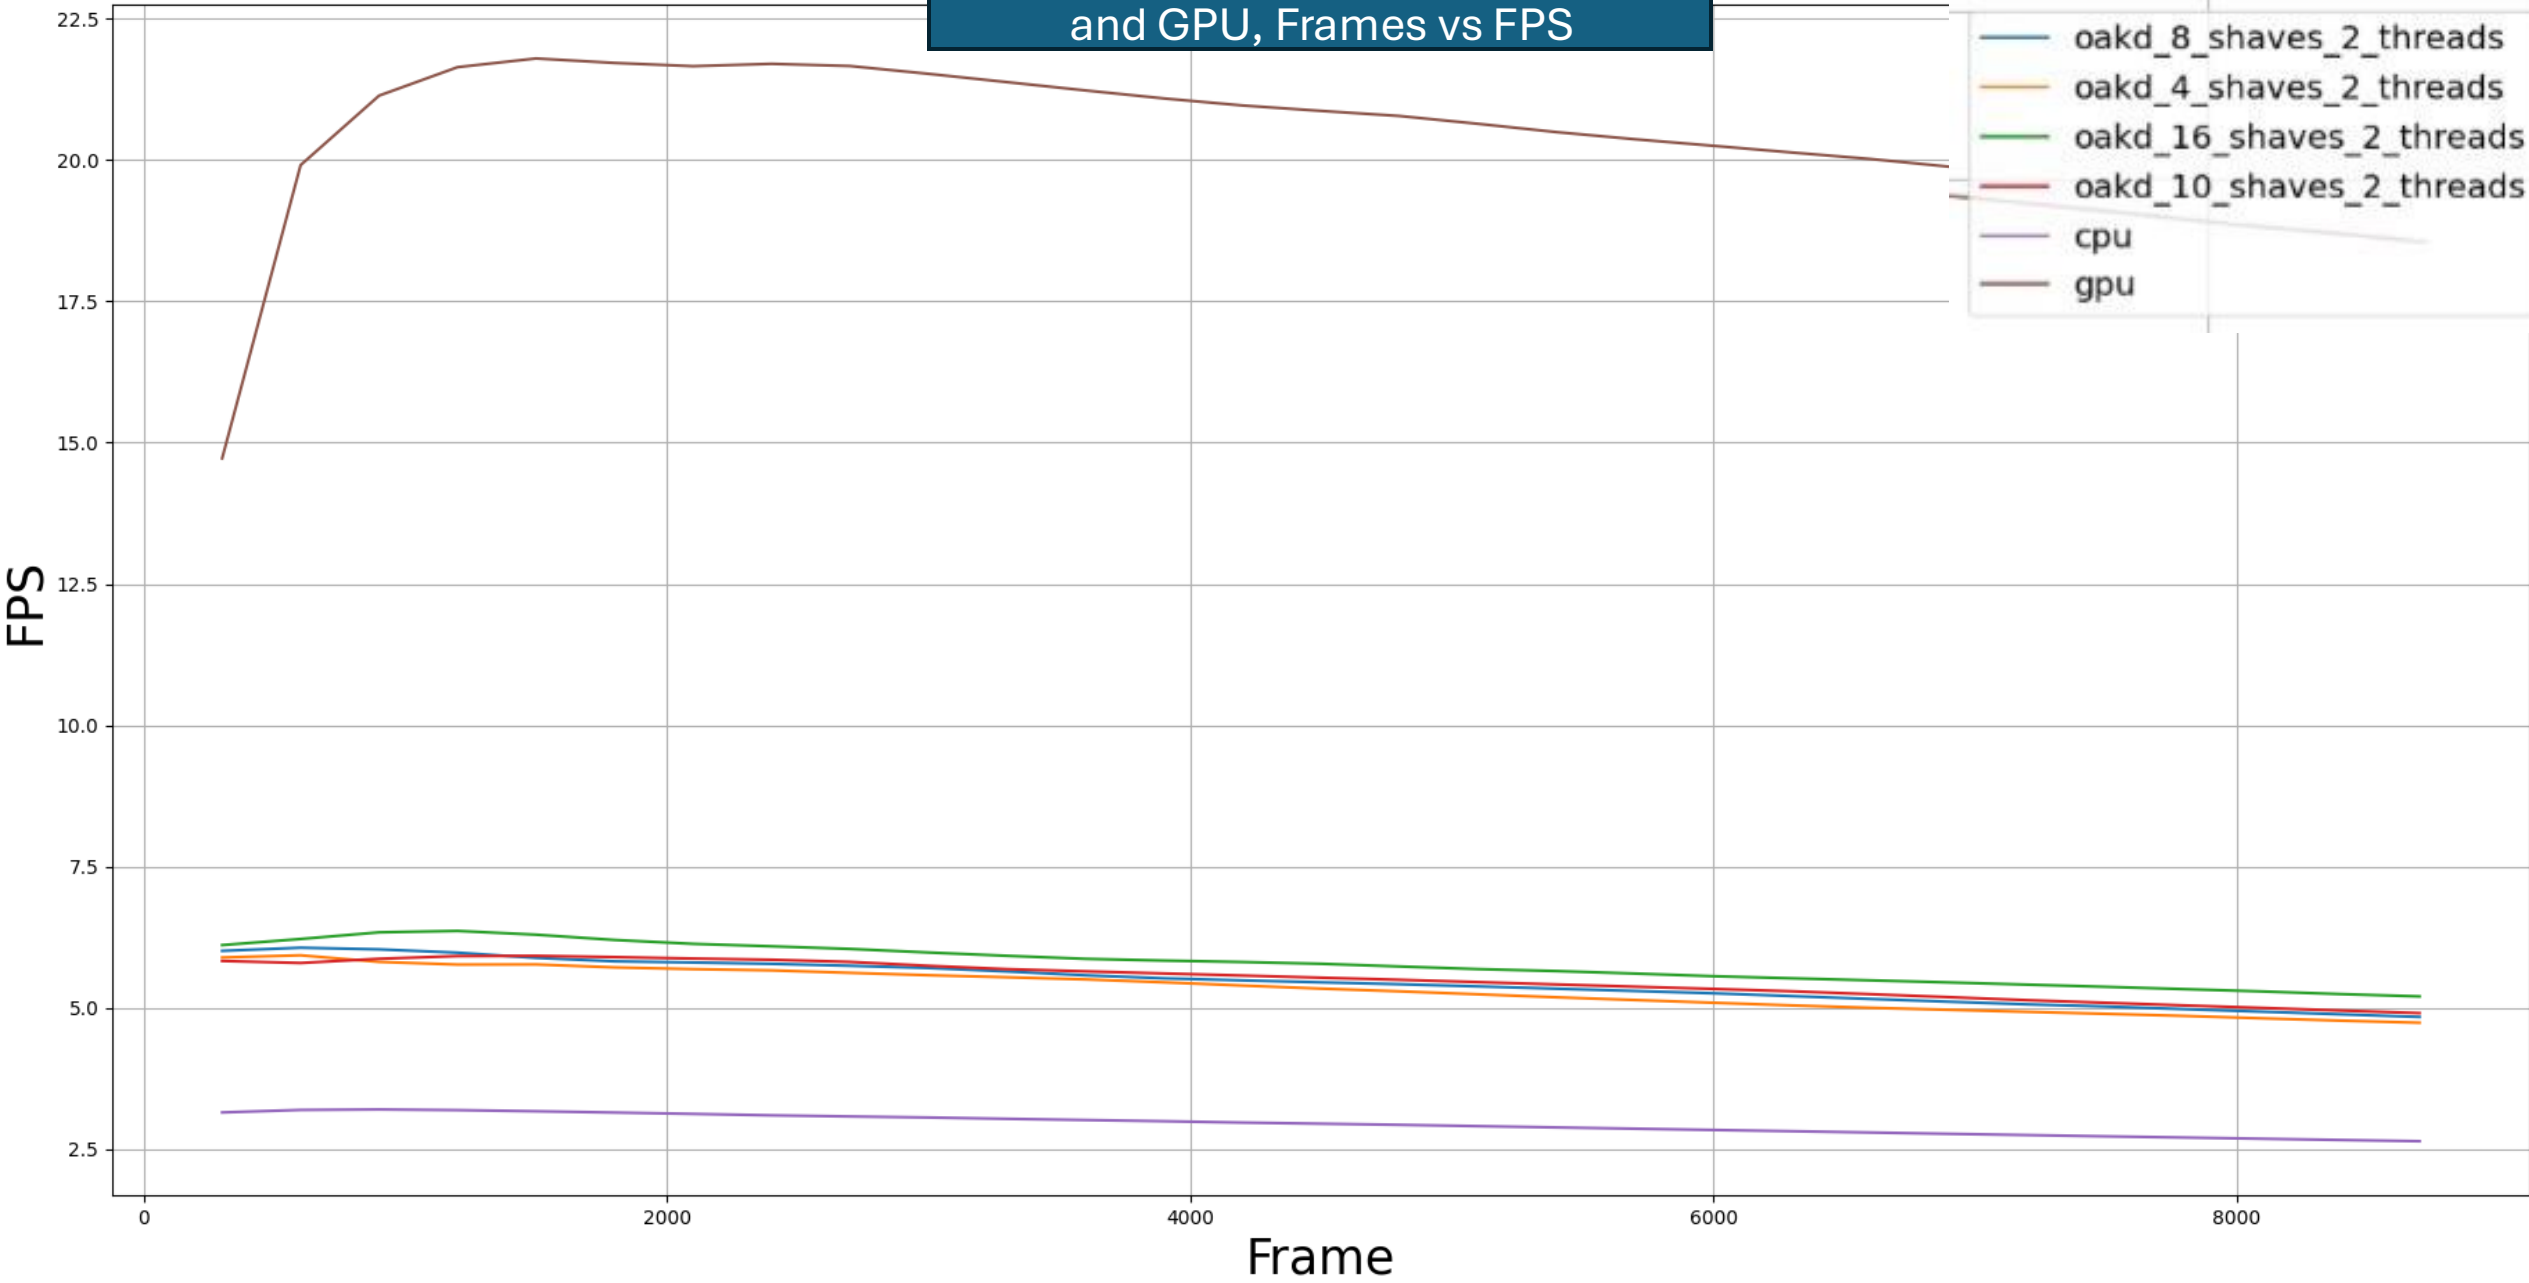
Power Consumption vs FPS per Benchmark Run




How # of Shaves Affects FPS




Same Oak-D Benchmarks with CPU
and GPU, Frames vs FPS



Frame vs Power Draw




Average Power Draw for Different Shaves



Same # of Shaves,
Different Threads



2 and 1 threads
compete for best fps




FPS

Same # of Shaves,
Different Threads

- oakd_10_shaves_3_threads
- oakd_10_shaves_2_threads
- oakd_10_shaves_1_threads

2 and 3 threads
compete for best fps.
Hard to make any
comparison on how
FPS changes with
different threads

Frame



Project Learnings

—

- Increasing number of shaves increased power consumption
- Increasing number of shaves increased inferencing frame rate
- Oak-D outperforms CPU in terms of performance and lower power consumption
- GPU outperforms Oak-D on FPS, but takes much more power
- Oak-D had the best power to FPS Ratio

Hurdles and Future Work

Project Hurdles



- Power Consumption on a Raspberry PI
 - Switched to Orin
- Testing Models with Different Shave Counts
 - You have to recompile the model for each one
- Automating Data Collection
 - Long inference times on videos to get accurate results

Future Work



- Try different file sizes/video qualities
- Inference directly from camera
- Try different models

Questions?